

Multithreading

VGP131-OOP II

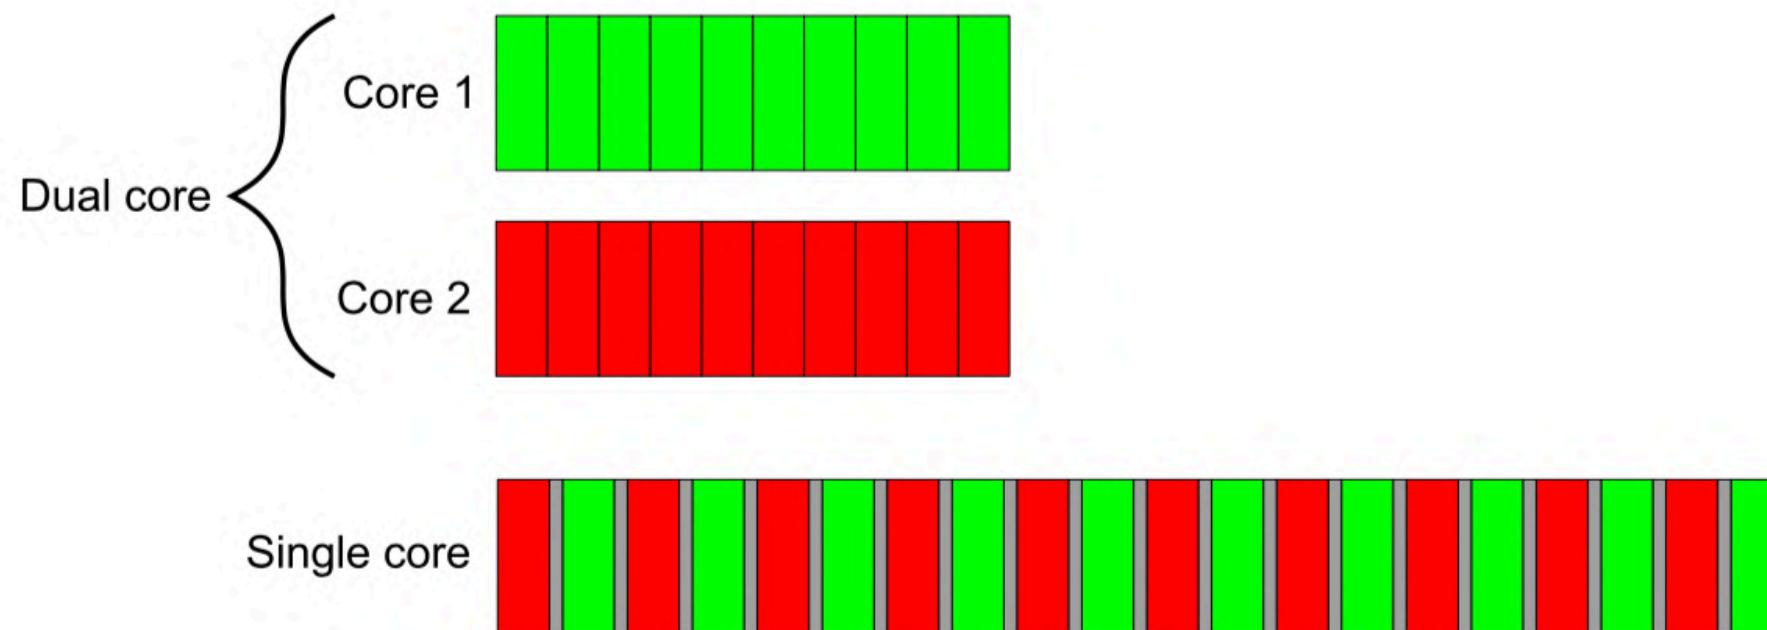
Instructor: Ivaldo Tributino

Concurrency in computer systems

- Concurrency in terms of computers means a single system performing multiple independent activities in parallel, rather than sequentially, or one after the other.
- Task switching: Computers with a single processing unit or core, can really only perform one task at a time, but it can switch between tasks many times per second that appear to be happening simultaneously.
- Whether they have multiple processors or multiple cores within a processor (or both), these computers are capable of genuinely running more than one task in parallel. We call this *hardware concurrency*.

Hardware Threads

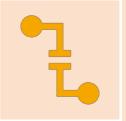
- The figure shows an idealized scenario of a computer with precisely two tasks to do. On a dual-core machine, each task can execute on its own core. On a single-core machine doing task switching, the chunks from each task are interleaved.



CONCURRENCY WITH MULTIPLE THREADS



Threads are much like lightweight processes: each thread runs independently of the others, and each thread may run a different sequence of instructions.



All threads in a process share the same address space, and most of the data can be accessed directly from all threads – global variables remain global, and pointers or references to objects or data can be passed around among threads.



Although it's often possible to share memory among processes, this is complicated to set up and often hard to manage, because memory addresses of the same data aren't necessarily the same in different processes.

Getting started

Every C++ program has at least one thread, which is started by the C++ runtime: the thread running `main()`.

Imperative that you ensure that the thread is correctly joined or detached, even in the presence of exceptions.

If you need to wait for a thread to complete, you can do this by calling `join()` on the associated `std::thread` instance.

Once you've called `join()`, the `std::thread` object is no longer joinable, and `joinable()` will return false.

Creating Threads

The screenshot shows a Visual Studio Code (VS Code) interface. On the left is a vertical sidebar with various icons for file management, search, and other tools. The main area displays a C++ code editor with the following content:

```
libraries.h      main.cpp  ●  Makefile
C main.cpp > function1()

12 void function1()
13 {
14     for(unsigned i = 0; i<5; i++){
15         for(unsigned j=0; j<5; j++){
16             cout << '1';
17         }
18         cout << '\n' ;
19     }
20 }

21

22 void function2()
23 {
24     for(unsigned i = 0; i<5; i++){
25         for(unsigned j=0; j<5; j++){
26             cout << '2';
27         }
28         cout << '\n' ;
29     }
30 }

31
32 int main(){
33
34     //***** ---- Creating Threads ---- *****
35     //----- Creating Threads -----
36     //*****
37
38     function1();
39     function2();
}

The function1( ) will be run first and then function2( ).
```

The terminal at the bottom shows the output of running the program:

```
(base) Ivaldo:Week10 admin$ ./main
11111
11111
11111
11111
11111
22222
22222
22222
22222
(base) Ivaldo:Week10 admin$
```

The status bar at the bottom indicates the file is on master*, has 0 changes, and is in Live Share mode.

Creating Threads

The screenshot shows a Microsoft Visual Studio Code (VS Code) interface. On the left, there's a vertical toolbar with various icons for file operations, search, and other development tools. The main workspace displays a C++ file named 'main.cpp' with the following code:

```
107
108 int main(){
109
110     //***** ---- Creating Threads ---- *****
111
112     // // function1();
113     // // function2();
114
115
116
117
118     std::thread t1(function1);
119     std::thread t2(function2);
120
121     t1.join();
122     t2.join();
123
124     if(!t1.joinable()){
125         cout << "thread object is no longer joinable" << endl;
126     }
127
128
129
130
131
132
133
134
```

Below the code editor, there are tabs for PROBLEMS, OUTPUT, TERMINAL, and DEBUG CONSOLE. The TERMINAL tab is active, showing the command-line output of the program:

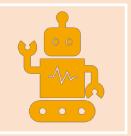
```
2222
111
11111
11111
11111
11111
11111
2
22222
22222
22222
thread object is no longer joinable
(base) Ivaldo:Week10 admin$
```

The status bar at the bottom indicates the current position is Line 123, Column 1, with 4 spaces, using UTF-8 encoding, and the file is saved in LF format (C++). There are also icons for navigating between files and live sharing.

Initializing thread with an object



In the previous example, we used regular function for the thread task. However, we can use any callable object such as lambda functions, functor and function pointer.



A functor is an object of a class that overloads operator () – function call operator.

Passing arguments to a thread function

- Passing arguments to the callable object or function is fundamentally as simple as passing additional arguments to the `std::thread` constructor. But it's important to bear in mind that by default the arguments are *copied* into internal storage, where they can be accessed by the newly created thread of execution, even if the corresponding parameter in the function is expecting a reference.
- To pass the message by reference, we must use **`std::ref(data)`** – A reference to the data will be passed instead of a reference to a copy of the data.
- We can use **`move`** to pass the parameter without copying and not sharing memory between the threads.

Using Functors

The screenshot shows a code editor interface with the following details:

- File Explorer:** Shows files: main.cpp, Polygon.h, Polygon.cpp, thread_guard.h, thread_safe_stack.h, and library.h.
- Code Editor:** The main window displays the `main.cpp` file. The code demonstrates two ways to initialize a thread with a functor object:
 - Using `std::thread` with a regular object:

```
cout << "--- Initializing thread with an object ---" << '\n';
{
    Polygon square(4);
    std::thread t3(square, 8);

    if(t3.joinable()) t3.join();
}
cout << "- Using std::ref ---" << '\n';
```

 - Using `std::thread` with a `std::ref` object:

```
{
    Polygon square(4);
    std::thread t3(std::ref(square), 8);

    if(t3.joinable()) t3.join();
}
```
- Terminal:** The terminal output shows the execution of the code. It prints:

```
--- Initializing thread with an object ---
Constructor Invoked
Copy Constructor Invoked
Copy Constructor Invoked
Polygon destroyed
Polygon destroyed
Polygon destroyed
- Using std::ref ---
Constructor Invoked
Polygon destroyed
```
- Bottom Status Bar:** Shows the current branch is `master*`, the file is `main.cpp`, the line is 70, column is 30, and the encoding is UTF-8.

Using RAII to wait for a thread to complete

- Resource Acquisition Is Initialization (RAII), holding a resource is a class invariant, and is tied to object lifetime: resource allocation (or acquisition) is done during object creation (specifically initialization), by the constructor, while resource deallocation (release) is done during object destruction (specifically finalization), by the destructor. In other words, resource acquisition must succeed for initialization to succeed.
- RAII idiom and provide a class that does the `join()` in its destructor, as in the following listing.

Thread Guard

The screenshot shows a code editor interface with a sidebar containing various icons and a vertical orange bar on the left labeled "Thread Guard". The main area displays the content of `thread_guard.h`. The code defines a class `thread_guard` that wraps a `std::thread` object. It includes an explicit constructor that takes a `std::thread` reference and a destructor that prints a message if the thread is joinable. It also includes a copy assignment operator that deletes the original object.

```
libraries.h    main.cpp    thread_guard.h    Makefile
C thread_guard.h > ...
1
2 #include "libraries.h"
3
4 class thread_guard
5 {
6     std::thread& _t;
7
8 public:
9     //Explicit, that is, it cannot be used for implicit conversions and copy-initialization.
10    explicit thread_guard(std::thread& t): _t(t) {}
11
12 ~thread_guard()
13 {
14     if(_t.joinable()) _t.join();
15     cout << "thread object was destroyed." << endl;
16 }
17 // to ensure that they're not automatically provided by the compiler.
18 thread_guard(thread_guard const&)=delete;
19
20 thread_guard& operator=(thread_guard const&)=delete;
21
22 };
```

PROBLEMS OUTPUT TERMINAL DEBUG CONSOLE

11111
thread object is no longer joinable
(base) Ivaldo:Week10 admin\$

bash bash

master*+ 0 △ 0 Live Share

Ln 22, Col 3 Spaces: 4 UTF-8 LF C++ Mac

REF & RAI

The screenshot shows a code editor interface with a vertical toolbar on the left containing icons for file operations, search, and other development tools. The main area displays a C++ file named `main.cpp`. The code demonstrates RAII and thread management:

```
main.cpp > main()
124     cout << "thread object is no longer joinable" << endl;
125 }
126
127 //*****
128 // - Initializing thread with an object -
129 //*****
130 cout << "--- Initializing thread with an object ---" << endl;
131
132 {
133     Polygon square(4);
134     std::thread t3(square, 8);
135
136     if(t3.joinable()) t3.join();
137 }
138
139
140
141 cout << "--- USING REF AND RAI ---" << endl;
142 {
143     Polygon square(4);
144     {
145         std::thread t3(std::ref(square), 8);
146         thread_guard tg(t3);
147     }
148     cout << "-----" << endl;
149 }
150
151
152
153
154
155
156
157
```

The code includes several `cout` statements to output messages. The `thread_guard` class is used to ensure the thread object is destroyed when it goes out of scope. The output window shows the following results:

```
--- USING REF AND RAI ---
Constructor Invoked
thread object was destroyed.
-----
Polygon destroyed
(base) Ivaldo:Week10 admin$
```

The status bar at the bottom indicates the current line (Ln 146), column (Col 28), and encoding (UTF-8). The interface also features a navigation bar with tabs for PROBLEMS, OUTPUT, TERMINAL, and DEBUG CONSOLE.

Transferring ownership of a thread

- We transfer the ownership of the thread by moving it.
- Putting `std::thread` objects in a `std::vector` is a step toward automating the management of those threads: rather than creating separate variables for those threads and joining with them directly, they can be treated as a group.

Std :: move

The screenshot shows a code editor interface with a dark theme. On the left is a vertical toolbar with various icons. The main area displays a C++ file named `main.cpp`. The code demonstrates thread ownership transfer and destruction. It includes comments, cout statements, std::thread objects, and a `thread_guard` class.

```
86 //*****
87 // - Transferring ownership of a thread -
88 //*****
89 cout << "--- Transferring ownership of a thread ---" << '\n';
90 {
91     std::thread t_1(function1);
92     std::thread t_2 = std::move(t_1);
93     if(!t_1.joinable()){
94         cout << "thread object is no longer joinable" << '\n';
95     }
96     if(!t_2.joinable()){
97         cout << "thread object is no longer joinable" << '\n';
98     };
99     thread_guard g(t_2);
100 }
```

The terminal output window shows the execution of the code. It prints the introductory message, then indicates that both threads are no longer joinable. Finally, it shows the destruction of the thread object, which outputs '11111' five times before exiting.

PROBLEMS OUTPUT TERMINAL DEBUG CONSOLE

```
--- Transferring ownership of a thread ---
thread object is no longer joinable
11111
11111
11111
11111
11111
thread object was destroyed.
(base) Ivaldo:Week10 admin$
```

TERMINAL

- bash
- bash

Ln 107, Col 1 Spaces: 4 UTF-8 LF C++ Mac

How Many Threads Should We Create?

- One feature of the C++ Standard Library that helps here is `std::thread::hardware_concurrency()`. This function returns an indication of the number of threads that can truly run concurrently for a given execution of a program. On a multicore system it might be the number of CPU cores.

How many Threads?

The screenshot shows a code editor interface with the following details:

- File Explorer:** Shows files: main.cpp, Polygon.h, Polygon.cpp, thread_guard.h, thread_safe_stack.h, and library.h.
- Sidebar:** Includes icons for file operations, search, and navigation.
- Code Editor:** Displays the main.cpp file with the following content:

```
91     {
92         std::thread t_1(function1);
93         std::thread t_2 = std::move(t_1);
94         if(!t_1.joinable()){
95             cout << "thread object is no longer joinable" << '\n';
96         }
97         if(!t_2.joinable()){
98             cout << "thread object is no longer joinable" << '\n';
99         };
100        thread_guard g(t_2);
101    }
102
103 //*****
104 // - std::thread::hardware_concurrency() -
105 //*****
106 cout << "--- How Many Threads Should We Create? ---" << '\n';
107
108 {
109     unsigned int c = std::thread::hardware_concurrency();
110     cout << "Number of cores: " << c << '\n';
111 }
112
113
114
115
116
117
118
119
120
121
```
- Terminal:** Shows the output of the program:

```
11111
11111
11111
11111
thread object was destroyed.
--- How Many Threads Should We Create? ---
Number of cores: 4
(base) Ivaldo:Week10 admin$
```
- Bottom Bar:** Includes tabs for PROBLEMS, OUTPUT, TERMINAL, DEBUG CONSOLE, and a terminal icon with a '+' sign.
- Right Panel:** Shows a large sidebar with multiple tabs and sections, likely a file browser or project manager.

Thread id

- Every thread has its unique identifier;
- Object of type `std::thread::id`;
- We obtained id calling the `get_id()`;
- Provide a total ordering for all distinct values. This allows them to be used as keys in associative containers, or sorted, or compared in any other way that you as a programmer may see fit;
- The Standard Library also provides `std::hash<std::thread::id>` so that values of type `std::thread::id` can be used as keys in the new unordered associative containers too.

Threads ID?

The screenshot shows a code editor interface with a dark theme. On the left, there is a vertical toolbar with various icons: a file icon, a search icon, a '10K+' icon, a play/pause icon, a copy/paste icon, a GitHub icon, a refresh/circular arrow icon, a question mark icon with a '1' notification, and a gear settings icon.

The main area displays a C++ program named `main.cpp`. The code creates three threads (`t_1`, `t_2`, `t_3`) and prints their IDs. It also checks if each thread is joinable and terminates them if so. The code uses `std::atomic` to synchronize access to the counter variable.

```
110 //*****  
111 // --- Thread ID ---  
112 //*****  
113 cout << "--- Thread ID ---" << '\n';  
114 using SC = std::chrono::steady_clock;  
115 auto dealine = SC::now() + std::chrono::milliseconds(1);  
116  
117 // Accesses to atomic objects can establish synchronization between threads.  
118 std::atomic<int> counter = 0;  
119  
120 //create 3 different threads  
121 std::thread t_1([&](){  
122     while(SC::now()<dealine)  
123         printf("t_1: %d\n", ++ counter);  
124  
125 });  
126 std::thread t_2([&](){  
127     while(SC::now()<dealine)  
128         printf("t_2: %d\n", ++ counter);  
129  
130 });  
131 std::thread t_3([&](){  
132     while(SC::now()<dealine)  
133         printf("t_3: %d\n", ++ counter);  
134  
135 });  
136 //get id of threads t_1, t_2 and t_3  
137 std::thread::id id1 = t_1.get_id();  
138 std::thread::id id2 = t_2.get_id();  
139 std::thread::id id3 = t_3.get_id();  
140  
141 if(t_1.joinable())  
142 {  
143     t_1.join();  
144     cout << "Thread with id " << id1 << " is terminated" << '\n';  
145 }  
146 if (t_2.joinable())  
147 {  
148     t_2.join();  
149     cout << "Thread with id " << id2 << " is terminated" << '\n';  
150 }
```

The status bar at the bottom shows the file name `master*`, line `Ln 149, Col 70`, and other system information like `Spaces: 4`, `UTF-8`, `LF`, `C++`, and `Mac`.

Threads ID?



PROBLEMS

OUTPUT

TERMINAL

DEBUG CONSOLE

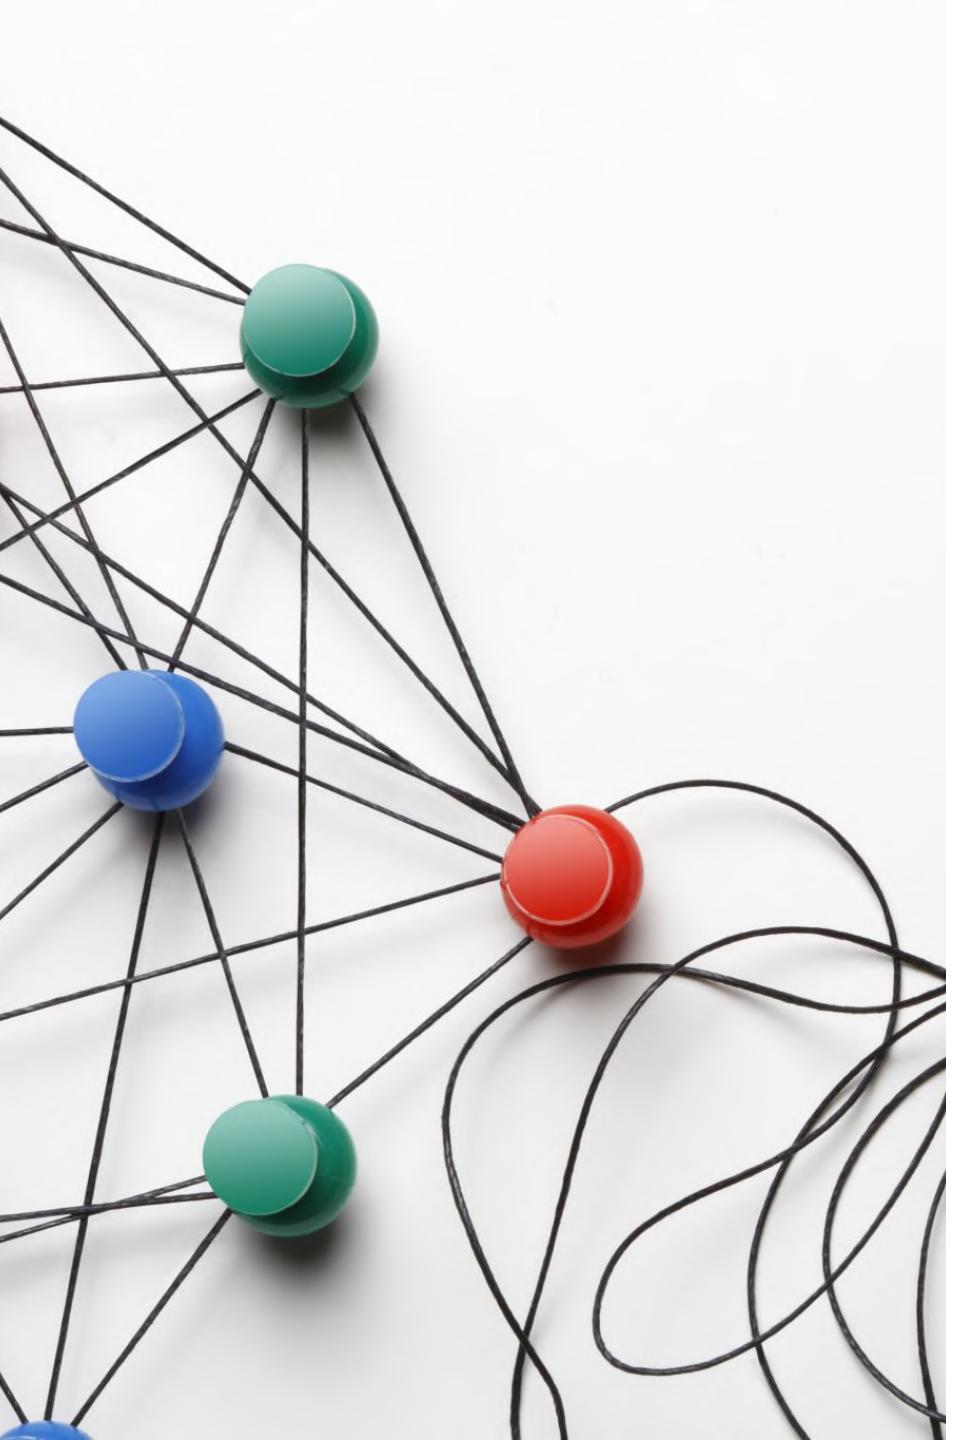
bash + ▾ ▷ ⌂ ✎ ×

```
t_2: 247
t_2: 248
t_2: 249
t_2: 250
t_2: 251
t_2: 252
t_2: 253
t_2: 254
t_2: 255
t_2: 256
t_2: 257
t_2: 258
t_2: 259
t_2: 260
t_2: 261
t_2: 262
t_2: 263
t_2: 264
t_2: 265
t_2: 266
t_2: 267
t_2: 268
t_2: 269
t_2: 270
t_2: 271
t_2: 272
t_2: 273
t_2: 274
t_2: 275
t_2: 276
t_2: 277
t_2: 278
t_2: 279
t_2: 280
t_2: 281
t_2: 282
t_2: 283
t_2: 284
t_2: 285
t_2: 286
t_2: 287
t_2: 288
t_2: 289
t_2: 290
t_2: 291
t_3: 185
t_1: 78
Thread with id 0x70000808d000 is terminated
Thread with id 0x700008110000 is terminated
Thread with id 0x700008193000 is terminated
Main thread id is :0x1178db5c0
(base) Ivaldo:Week10 admin$ █
```



✖ ↻ master*+ ⏪ ⏴ 0 △ 0 🔍 Live Share

Ln 125, Col 8 Spaces: 4 UTF-8 LF C++ Mac 🔍



Sharing data between threads

- If you're sharing data between threads, you need to have rules for which thread can access which bit of data when, and how any updates are communicated to the other threads that care about that data.
- Incorrect use of shared data is one of the biggest causes of concurrency-related bugs.
- The most basic mechanism for protecting shared data provided by the C++ Standard is the *mutex*.

Using mutexes in C++

- In C++, you create a mutex by constructing an instance of `std::mutex`, lock it with a call to the member function `lock()`, and unlock it with a call to the member function `unlock()`.
- The Standard C++ Library provides the `std::lock_guard` class template, which implements that RAII idiom for a mutex; it locks the supplied mutex on construction and unlocks it on destruction, thus ensuring a locked mutex is always correctly unlocked.

Mutex

The screenshot shows the VS Code interface with the following details:

- File Explorer:** Shows files: main.cpp, Polygon.h, Polygon.cpp, and thread_guard.h.
- Search Bar:** Shows "main()".
- Code Editor:** Displays the main.cpp file with the following code:

```
144 //*****
145 //      -- Using mutexes in C++ --
146 //*****
147 cout << "--- Using mutexes in C++ ---" << '\n';
148
149 vector<int> v;
150
151 std::thread thr_push(push,std::ref(v));
152 std::thread thr_pop(pop,std::ref(v));
153
154 if (thr_push.joinable()) thr_push.join();
155 if (thr_pop.joinable()) thr_pop.join();
156
157
158
159
160
161
162
163
164
```
- Terminal:** Shows the output of the program:

```
--- Using mutexes in C++ ---
Push 0
Push 1
Push 2
Push 2
Pop 1
Push 3
Push 4
Push 5
Pop 4
Push 6
Push 7
Push 8
Pop 7
Push 9
Pop 9
Pop 8
Pop 6
Pop 5
Pop 3
Pop 2
Pop 2
```
- Bottom Status Bar:** Shows "Ivaldo:Week10 admin\$".

```
std::mutex m;
void push(vector<int>& v)
{
    // m.lock();
    for (int i = 0; i < 10; ++i)
    {
        cout << "Push " << i << '\n';
        std::this_thread::sleep_for(std::chrono::seconds(1));
        v.push_back(i);
    }
    // m.unlock();
}

void pop(vector<int>& v)
{
    // m.lock();
    for (int i = 0; i < 10; ++i)
    {
        if(v.size()!=0)
        {
            int val = v.back();
            v.pop_back();
            cout << "Pop " << val << '\n';
        }
        std::this_thread::sleep_for(std::chrono::seconds(3));
    }
    // m.unlock();
}
```

Mutex

The screenshot shows a C++ development environment with several tabs at the top: main.cpp, Polygon.h, Polygon.cpp, and thread_g... . The main.cpp tab is active, displaying code related to mutexes and threads. The code includes a section for pushing values into a vector and another for popping them. The output terminal shows the execution of the program, which prints "Push" and "Pop" operations for indices 0 through 9.

```
void push(vector<int>& v)
{
    m.lock();
    for (int i = 0; i < 10; ++i)
    {
        cout << "Push " << i << '\n';
        std::this_thread::sleep_for(std::chrono::seconds(1));
        v.push_back(i);
    }
    m.unlock();
}

void pop(vector<int>& v)
{
    m.lock();
    for (int i = 0; i < 10; ++i)
    {
        if(v.size()!=0)
        {
            int val = v.back();
            v.pop_back();
            cout << "Pop " << val << '\n';
        }
        std::this_thread::sleep_for(std::chrono::seconds(3));
    }
    m.unlock();
}
```

PROBLEMS OUTPUT TERMINAL DEBUG CONSOLE

```
--- Using mutexes in C++ ---
Push 0
Push 1
Push 2
Push 3
Push 4
Push 5
Push 6
Push 7
Push 8
Push 9
Pop 9
Pop 8
Pop 7
Pop 6
Pop 5
Pop 4
Pop 3
Pop 2
Pop 1
Pop 0
(base) Ivaldo:Week10 admin$
```

Ln 312, Col 1 Spaces: 4 UTF-8 LF C++ Mac 🔍 🔍

Mutex

Waiting for an event or other condition

If one thread is waiting for a second thread to complete the task, it has several options:

- First, it could just keep checking a flag in shared data (protected by a mutex) and have the second thread set the flag when it completes the task.
- A second option is to have the waiting thread sleep for small periods between the checks using the `std::this_thread::sleep_for()`.
- The third, and preferred, option is to use the facilities from the C++ Standard Library to wait for the event itself. The most basic mechanism for waiting for an event to be triggered by another thread - `std::condition_variable`.

Waiting for an event or other Condition

The screenshot shows a Visual Studio Code (VS Code) interface with the following details:

- File Explorer:** On the left, there are icons for files like main.cpp, Polygon.h, Polygon.cpp, thread_guard.h, thread_safe_stack.h, and library.h. A "10K+" icon is visible.
- Code Editor:** The main area displays C++ code for waiting on a condition variable. The code includes comments explaining the steps: waiting for an event, printing the current thread ID, getting the current time, and sleeping until the next minute. It then sleeps for 5 seconds and prints the current time again.
- Terminal:** At the bottom, the terminal window shows the output of the program's execution. The output includes:

```
--- Waiting for an event or other condition ---
The id of current thread is 0x11171d5c0
Current time: 12:54:28
Waiting for the next minute to begin...
12:55:00 reached!
Current time: 12:55:05
(base) Ivaldo:Week10 admin$
```
- Bottom Status Bar:** Shows file path (master*+), line count (0 △ 0), Live Share, Ln 158, Col 1 (349 selected), Spaces: 4, UTF-8, LF, C++, Mac, and a few small icons.

Waiting for an event or other condition

The screenshot shows a Visual Studio Code (VS Code) interface with the following details:

- File Explorer:** Shows files: main.cpp (1), Makefile, library.h, Polygon.h, and Polygon.cpp.
- Sidebar:** Includes icons for file operations, search, and other development tools.
- Code Editor:** Displays the main.cpp file with the following code:

```
216 cout << "--- By condition " << '\n';
217
218 std::thread thr_print(print_name);
219 std::thread thr_input(input_name);
220
221 if(thr_print.joinable()) thr_print.join();
222 if(thr_input.joinable()) thr_input.join();
223
224
225 void print_name () {
226     std::unique_lock<std::mutex> lck(m);
227     while (!ready) cv.wait(lck);
228     string name;
229     std::cin >> name;
230     std::cout << "My name is: " << name << '\n';
231 }
232
233 void input_name() {
234     std::unique_lock<std::mutex> lck(m);
235     cout << "Enter your name: ";
236     ready = true;
237     cv.notify_one();
238 }
```

A yellow box highlights the sections of code within the `print_name()` and `input_name()` functions.
- Terminal:** Shows the output of the program:

```
--- By condition
Enter your name: Ivaldo
My name is: Ivaldo
(base) Ivaldo:Week10 admin$
```
- Bottom Bar:** Includes tabs for PROBLEMS, OUTPUT, TERMINAL (selected), and DEBUG CONSOLE, along with icons for bash, terminal, and file operations.

Asynchronous task: std::future

Tasks work very similar to threads, but the main difference is that they can return a value.

You use `std::async` to start an *asynchronous task* for which you don't need the result right away.

`std::async` returns a `std::future` object, which will eventually hold the return value of the function.

When you need the value, you just call `get()` on the future, and the thread blocks until the future is *ready* and then returns the value.

`std::async` function template declared in the `<future>` header.

$$\begin{bmatrix} A_x & A_y & A_z \end{bmatrix} \begin{bmatrix} B_x \\ B_y \\ B_z \end{bmatrix} = A_x B_x + A_y B_y + A_z B_z = \vec{A} \cdot \vec{B}$$

```
long long getDotProduct(std::vector<int>& v, std::vector<int>& w){  
    return std::inner_product(v.begin(), v.end(), w.begin(), 0LL);  
}  
  
long long getDotProductFuture(std::vector<int>& v, std::vector<int>& w){  
  
    auto future1 = std::async([&]{return std::inner_product(&v[0], &v[v.size()/4], &w[0], 0LL);});  
    auto future2 = std::async([&]{return std::inner_product(&v[v.size()/4], &v[v.size()/2], &w[v.size()/4], 0LL);});  
    auto future3 = std::async([&]{return std::inner_product(&v[v.size()/2], &v[v.size()*3/4], &w[v.size()/2], 0LL);});  
    auto future4 = std::async([&]{return std::inner_product(&v[v.size()*3/4], &v[v.size()], &w[v.size()*3/4], 0LL);});  
  
    return future1.get() + future2.get() + future3.get() + future4.get();  
}
```

Using std::future

The screenshot shows a code editor interface with several tabs at the top: main.cpp, Polygon.h, Polygon.cpp, thread_guard.h, thread_safe_stack.h, and library.h. The main.cpp tab is active. A search bar indicates the code contains 'NUM'.

The code in main.cpp includes the following sections:

- Random number generation and distribution setup.
- Vector creation and filling.
- Measurement of execution time for the standard dot product calculation.
- A highlighted section containing two functions: `getDotProduct` and `getDotProductFuture`.
- Measurement of execution time for the future-based dot product calculation.
- A return statement at the end.

The highlighted code block contains:

```
long long getDotProduct(std::vector<int>& v, std::vector<int>& w){  
    return std::inner_product(v.begin(), v.end(), w.begin(), 0LL);  
}  
  
long long getDotProductFuture(std::vector<int>& v, std::vector<int>& w){  
  
    auto future1 = std::async([&]{return std::inner_product(&v[0], &v[v.size()/4], &w[0], 0LL);});  
    auto future2 = std::async([&]{return std::inner_product(&v[v.size()/4], &v[v.size()/2], &w[v.size()/4], 0LL);});  
    auto future3 = std::async([&]{return std::inner_product(&v[v.size()/2], &v[v.size()*3/4], &w[v.size()/2], 0LL);});  
    auto future4 = std::async([&]{return std::inner_product(&v[v.size()*3/4], &v[v.size()], &w[v.size()*3/4], 0LL);});  
  
    return future1.get() + future2.get() + future3.get() + future4.get();  
}
```

At the bottom, the terminal output shows the results of running the code:

```
--- Using std::future ---  
getDotProd(v,w): 250004960777  
Total Time Taken = 1.51819  
getDotProdFuture(v,w): 250004960777  
Total Time Taken using Future = 0.141271  
(base) Ivaldo:Week10 admin$
```

The status bar at the bottom right shows: Ln 17, Col 1 (33 selected) Spaces: 4 UTF-8 LF C++ Mac.

References:

1. Bogotobogo (2022/03/02) Thread Tutorials, https://www.bogotobogo.com/cplusplus/C11/1_C11_creating_thread.php
2. CS & Math Stuff(2022/03/03) Threads-Basics, <https://oneraynyday.github.io/dev/2017/11/19/C++-Threads-Basics/>
3. Hackernoon (2022/03/02) Learn C++ Multi-Threading in 5 Minutes, <https://hackernoon.com/learn-c-multi-threading-in-5-minutes-8b881c92941f>
4. Tutorialcup (2022/03/02) C++ Multithreading , <https://www.tutorialcup.com/cplusplus/multithreading.htm>
5. Baeldung (2022/03/02) The Difference Between Asynchronous And Multi-Threading,
<https://www.baeldung.com/cs/async-vs-multi-threading>
6. MC++ (2022/03/02) Asynchronous Function Calls, <https://www.modernescpp.com/index.php/asynchronous-function-calls>
7. Anthony, Williams. "c++ concurrency in Action. Practical Multithreading." (2001).
8. Youtube, Back to Basics: Concurrency - Arthur O'Dwyer - CppCon 2020.
<https://www.youtube.com/watch?v=F6lpn7gCOsY>

Appendix

- ❖ Libraries.h
- ❖ thread_guard.h
- ❖ Polygon.h
- ❖ Polygon.cpp
- ❖ main.cpp